WorkRadar

Product Requirements Document

Author · Eswar Mahalingam
Version 1.0 · 12 Oct 2026
Status · MVP built on mock data

THE FOUR QUESTIONS

1. **Top work for the day.** What is the team carrying, in rank order?
2. **What happened today.** What changed since I last looked?
3. **What needs immediate attention.** Where does my action change an outcome?
4. **Who is blocked.** Who is waiting, on what, and for how long?

POSITIONING

For managers who lose the first part of every day to status-hunting, WorkRadar is a **daily attention list**. Unlike a dashboard per tool, it ranks across tools and explains every rank.

TARGET USERS

- **Primary:** first-line engineering and delivery managers with 5 to 12 direct reports who work in Slack, Jira, GitHub and email.
- **Secondary:** programme and delivery leads who track 2 to 4 squads.
- **Indirect:** team members whose blockers get seen without a status ping.
- **Not for:** personal task lists or executive portfolio reporting.

DESIGN PRINCIPLES

- **Summary before detail.** A four-sentence brief first, lists second, threads last.
- **Show the reason for every rank.** No black-box priority.
- **Every action can be undone.** Assign, Snooze and Mark done all offer Undo.
- **Fail honestly.** If a tool or the AI is down, say so and show the last good copy.

IN THIS PACKAGE

- PRD (this document)
- UX screens: 5 high-fidelity boards
- Working MVP and source code
- AI workflow document
- Demo video script

One screen that tells a manager what needs them today

WorkRadar reads Slack, Jira, GitHub and email, ranks every open item with a visible score, and answers a manager's four daily questions in under five minutes.

01 Problem statement

A manager of a software team starts the day by rebuilding status from four places: Slack threads, Jira boards, GitHub pull requests and email. Each tool shows its own slice. None ranks across tools, and none says who is stuck or for how long. The manager becomes the integration layer, like a kirana owner tallying four separate ledgers before opening the shutter. The cost shows up as blockers that age unseen until standup, reviews that wait a day, and decisions with a 5 PM deadline that surface at 4:45.

> **Job to be done.** When I start my day, I want to see where my attention changes an outcome, so that nobody on my team waits on me without my knowing.

02 User personas

PRIMARY

Priya Menon

Engineering Manager, Checkout squad, 6 engineers.
Goal: unblock people fast and protect release dates.
Frustration: hears about blockers at standup, a day late.
Needs: "Tell me what only I can unblock."

SECONDARY

Arjun Bhatt

Delivery Lead across 3 squads.
Goal: spot release risk a week early.
Frustration: status is only as good as the last Jira update.
Needs: one ranked risk list across teams and tools.

INDIRECT

Ananya Rao

Senior Backend Engineer.
Goal: raise a blocker once and get back to work.
Frustration: repeats the same status in three tools.
Needs: her wait time visible to her manager automatically.

03 User journey

STAGE	TODAY	WITH WORKRADAR	SCREEN
Start of day	Opens four tools and scrolls unread messages	Reads a four-sentence brief built from the ranked items	Dashboard
Triage	Ranks in the head; misses links across tools	"Needs you now" list with a visible score per item	Dashboard, Feed
Understand	Reads a 30-message thread for context	AI digest, suggested next step and "why it is ranked here"	Item details
Act	Switches to the source tool; follow-ups get lost	Assign, Snooze or Mark done in one click, with Undo	Feed, Details
Follow up	Re-polls every tool after lunch	"What happened today" and "Who is blocked" with wait time	Dashboard
Find	Searches each tool separately	One search across all tools, including done and snoozed	Search

04 Key pain points

- **Fragmentation.** Status lives in four tools and nothing ranks across them.
- **Invisible blockers.** Nobody tracks how long a person has been waiting.
- **Deadline blindness.** Decisions that arrive by email or DM carry due times no board shows.
- **Context switching.** Every tool change breaks focus and follow-ups slip.
- **Thread overload.** Context means reading the whole thread.
- **Ownerless risks.** Alerts and vendor notices sit unassigned.
- **Low trust in automation.** Managers ignore a priority they cannot explain.
- **Late escalation.** Engineers hesitate to escalate, so managers act only when asked.

These are hypotheses drawn from the brief. Validation plan: five manager interviews and a one-week diary of morning triage before Phase 1.

WorkRadar

Scope, assumptions and risks

Sections 05 to 08

MVP SCOPE · IN

- Dashboard that answers the four questions
- AI brief with a rule-based fallback
- Priority feed with score, filters and actions
- Item detail page and cross-tool search
- Empty and error states
- Mock data: 17 items, 4 tools, 6 teammates

MVP SCOPE · OUT

- Live connectors and sign-in
- Backend, database, multi-user state
- Writing changes back to Jira, GitHub or Slack
- Notifications and scheduled digests

ASSUMPTIONS

- Managers already have read access to all four tools for their team.
- Ticket keys and PR numbers are enough to link the same work across tools.
- A transparent rule-based score earns trust faster than a learned one at this stage.
- The employer approves the AI tool for work data before any real connector is switched on.
- Mock data is acceptable for this MVP, as the brief states.

RISKS AND MITIGATIONS

- **Wrong ranking hides a real fire.** Score is visible and every factor is listed; incidents carry the highest base weight.
- **AI states something false.** Brief uses only ranked item data and is labelled; fallback is rule-based.
- **Connector outage.** Cached copy is labelled with a retry.
- **Privacy.** DMs and email need explicit opt-in scopes.

DEPENDENCIES

- API access to Jira, GitHub, Slack and the mail provider
- An approved AI model for summaries

OPEN QUESTIONS

- Should Snooze hide an item from the whole team or only from me?

05 Success metrics

METRIC	DEFINITION	PILOT TARGET
North Star: triage time	Median time from first open of the day to the manager's first action	5 minutes or less
Activation	Managers who take at least one action from the feed on day 1	70% or more
Blocker age	Median hours a teammate stays blocked before the blocker is cleared	30% lower than week 0
Ranking trust	Share of the top 5 items the manager acts on or marks as correctly ranked	60% or more
Guardrail: AI accuracy	Briefs reported as containing a wrong fact	Under 2 per 100 briefs
Guardrail: missed fires	Critical items that were open but ranked outside the top 10	Zero

Targets are for a four-week pilot with five managers. No baseline exists yet; week 0 measures current triage time and blocker age.

06 Features: must-have and nice-to-have

MUST-HAVE · IN THE MVP

- **Dashboard:** brief, radar, Needs you now, Top work today, Who is blocked, What happened today
- **AI-generated summary** of the ranked items, with a rule-based fallback
- **Priority feed:** one ranked list, filters by status, lens and tool
- **Item details:** digest, next step, thread, score breakdown
- **Search** across all tools with match highlighting
- **Actions:** Assign, Mark done, Snooze, each with Undo
- **Empty and error states** for feed, search and sync failure

NICE-TO-HAVE · AFTER THE MVP

- Live read-only connectors: Jira and GitHub first, then Slack and email
- Write-back: reassign in Jira, approve in GitHub, reply in Slack
- 9 AM brief delivered to Slack or email
- "Ranked wrong" feedback that tunes weights per manager
- Team load view: open items and blocked hours per person
- Roll-up across squads for delivery leads
- Single sign-on, audit log and admin controls
- Calendar awareness: leave and release dates feed the score

07 How priority is calculated

FACTOR	POINTS	FACTOR	POINTS
Type: incident / blocker / decision	`58 / 34 / 28`	Blocks a person: over 24 h / under 24 h	`+14 / +9`
Type: review / risk / planned work / FYI	`22 / 20 / 12 / 4`	Customer impact stated	`+12`
Overdue / due today / due tomorrow	`+24 / +18 / +10`	Waiting on the manager; a release depends on it	`+8 each`
Due within three days	`+4`	No owner	`+6`

The score is the sum, capped at 99. **Critical** is 70 and above, **High** 48 to 69, **Medium** 28 to 47, **Low** below 28. Each item page lists its own factors.

08 Phased plan

PHASE	WHAT SHIPS	EXIT CHECK
0 · Done	Clickable MVP on mock data: front end only, no sign-in, no database	All five brief requirements demoed
1	Read-only Jira and GitHub connectors, sign-in, hosted database	Five pilot managers onboarded
2	Slack and email connectors, live AI digests, "ranked wrong" feedback	Triage time and trust targets met
3	Write-back actions, 9 AM digest, multi-squad roll-up	Security, scale and licence review passed

Non-goals. WorkRadar does not replace the four tools, does not score individual performance, and never acts in a source tool without the manager's click.